

# Introduction

Ravindra Kudache

# Who ?

## **G-Research**

*“G-Research is a leading quantitative research and technology company. By using the latest scientific techniques, we produce world-beating predictive research and build advanced technology to analyse the world’s data”.*

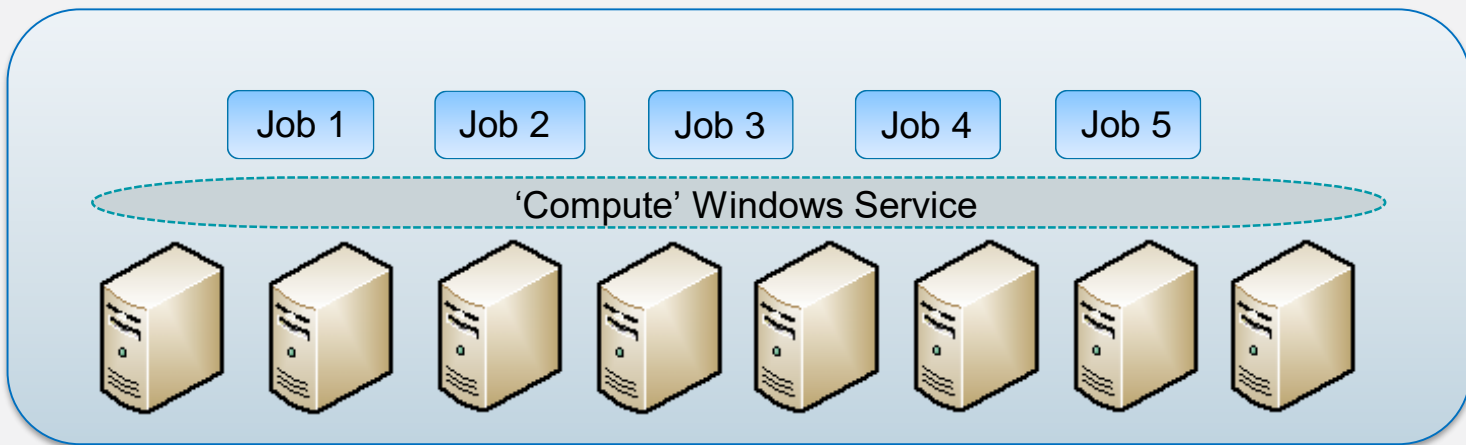
## **Windows Engineering**

Windows Engineering are responsible for automated Windows provisioning, software delivery and configuration management across the estate. We strongly adhere to the principles of IaC (Infrastructure-as-Code).

# What ?

These 1000 physical windows servers makeup a High-Throughput Computing System, which we refer to as the 'Farm'. We have multiple farms hosting tens of thousands of servers and this is just one small subset.

The Farm provides compute resource that allows for many computational tasks to be completed over a long period of time. Mechanisms in the farm include; job queueing, scheduling, prioritizing and resource management. Nodes are predominantly **'stateless'**




# Why ?

< 30 Days


- Our objective is to have no nodes in the farm older than **30 days**
- This prevents the build up of unintentional artefacts on older nodes
- This allows frequent updating of software
- Servers do not require patching
- The automated mechanism allows for more consistent node builds
- Reduce manual intervention and time spent on farm builds (5 teams)
- Reduce troubleshooting time.

# The Problem

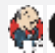


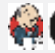
# The Rebuild Process

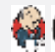
# Process Overview




# Process Overview


 Scheduler →

 Drainer →

 Rebuilder →

In Farm 

Out of Farm 





# Process Overview - Drainer

```
- hosts: lifecycle_scheduled
gather_facts: False
name: SCHEDULED - Kick off draining for scheduled nodes
serial: 1
vars:
  profile_live_count: "{{ (groups['lifecycle_live'] | intersect(groups[profile_group])) | length }}"
  profile_rebuilding_count: "{{ (groups['lifecycle_rebuilding'] | intersect(groups[profile_group])) | length }}"
  profile_draining_count: "{{ (groups['lifecycle_draining'] | intersect(groups[profile_group])) | length }}"
  profile_scheduled_count: "{{ (groups['lifecycle_scheduled'] | intersect(groups[profile_group])) | length }}"
tasks:
  - name: SCHEDULED - Debug Profile counts
    debug:
      msg:
        - "profile_group: {{ profile_group }}"
        - "live: {{ profile_live_count }}"
        - "rebuilding: {{ profile_rebuilding_count }}"
        - "draining: {{ profile_draining_count }}"
        - "scheduled: {{ profile_scheduled_count }}"
        - "rebuilding_queue: {{ rebuilding_queue }}"
  - name: SCHEDULED - Drain host
    include_tasks: tasks/drain_host.yml
    when:
      - ((profile_rebuilding_count|int + profile_draining_count|int) < rebuilding_queue)
      - ((profile_live_count|int + profile_scheduled_count|int - 1) > 0)
      - rebuild_supported
```

```
#Refresh the inventory as things have changed
- meta: refresh_inventory
```

# Process Overview

## **Live -> Scheduled**

- Used to see the overall health of the process
- All nodes ending up in scheduled indicates rate issues

## **Scheduled -> Draining**

- Rate limited by the number of nodes in Draining and Rebuilding
- Uses ansible inventory to calculate if transition is possible

## **Draining -> Rebuilding**

- Checks whether jobs are still running on the node
- Triggers external “Rebuilder” Jenkins job/playbook

## **Rebuilding -> Live**

- Processed by generic rebuild process on a fixed version
- Time taken only needs to be of the order of draining time

# Inventory Overview

	Overall Inventory			
	Live	Scheduled	Draining	Rebuilding
Profile A	1	1	1	1
Profile B	1	1	1	1
Profile C	1	1	1	1
Profile A	Profile B	Profile C	Profile D	
4	4	4	4	4

# Ansible Inventory V1

CMDB

Live	Scheduled	Draining	Rebuilding
4	4	4	4

# Ansible Inventory V1

## Inventory Files

```

  ▾ inventories
    ▾ farm_a
      ▾ group_vars
        ! all.yml
        ! farm_a.cmdb.yml
      ▾ farm_b
        ▾ group_vars
          ! all.yml
          ! farm_b.cmdb.yml

```

```

inventories > farm_a > ! farm_a.cmdb.yml
1  plugin: cmdb
2  query: farm eq farm_a

```

## Issues:

- All Compute-Farm facts had to be calculated per node post inventory
- Each nodes awareness of other nodes had to be calculated separately
- Calculation of compute farm groups was done in YAML 'code'
- Setting variables per 'profile' had to be done post inventory
- Playbooks were hard to follow

# Ansible Inventory V2

CMDB + Compute Farm

	Live	Scheduled	Draining	Rebuilding
Profile A	1	1	1	1

# Ansible Inventory V2

## Inventory Files

```
✓ inventories
  ✓ farm_a-profile_a
    ✓ group_vars
      ! all.yml
      ! farm_a-profile_a.compute.yml
    ✓ farm_a-profile_b
      ✓ group_vars
        ! all.yml
        ! farm_a-profile_b.compute.yml
  ✓ farm_b-profile_a
    ✓ group_vars
      ! all.yml
      ! farm_b-profile_a.compute.yml
  ✓ farm_b-profile_b
    ✓ group_vars
      ! all.yml
      ! farm_b-profile_b.compute.yml
```

```
inventories > farm_a-profile_a > ! farm_a-profile_a.yml
1  plugin: compute
2  cmdb_query: farm eq farm_a
3  compute_query: profile eq profile_a
```

## Advantages

- Simple to get number of nodes out of farm
- Possible to set per profile variables
- Ansible playbooks are much simpler
- Easier for external teams to contribute
- Able to use “refresh\_inventory” to update node states

## Issues:

- Each farm requires multiple playbooks runs via Jenkins to cover the whole farm
- Setting per general profile variables has to be done in every inventory
- Setting up support for new profiles is still complex

# Ansible Inventory V3

CMDB + Compute Farm

	Live	Scheduled	Draining	Rebuilding
Profile A	1	1	1	1
Profile B	1	1	1	1
Profile C	1	1	1	1
Profile D	1	1	1	1



# Ansible Inventory V3

## Inventory Files



```
inventories > farm_a > ! farm_a.compute.yml
```

```
1 plugin: compute
2 cmdb_query: farm eq farm_a
```

```
inventories > farm_a > group_vars > ! profile_a.yml
```

```
1 rebuild_supported: True
```

## Advantages

- Simple to add profiles
- Single playbook run per farm
- All nodes included even when not rebuild supported
- Even easier for external teams to contribute

## Issues:

- Getting the profile + lifecycle cross section is more complex:

```
profile_live_count: "{{ (groups['lifecycle_live'] | intersect(groups[profile_group])) | length }}"
profile_rebuilding_count: "{{ (groups['lifecycle_rebuilding'] | intersect(groups[profile_group])) | length }}"
profile_draining_count: "{{ (groups['lifecycle_draining'] | intersect(groups[profile_group])) | length }}"
profile_scheduled_count: "{{ (groups['lifecycle_scheduled'] | intersect(groups[profile_group])) | length }}"
```

# Logging

## **Ansible:**

- Logs playbook, play, and task results to splunk via a callback plugin
- Includes timing information
- Sends chat messages on error

## **Jenkins:**

- Full playbook output retained for 30 days
- Job meta information logged to splunk

## **Compute Farm:**

- Snapshot of node states are logged to splunk

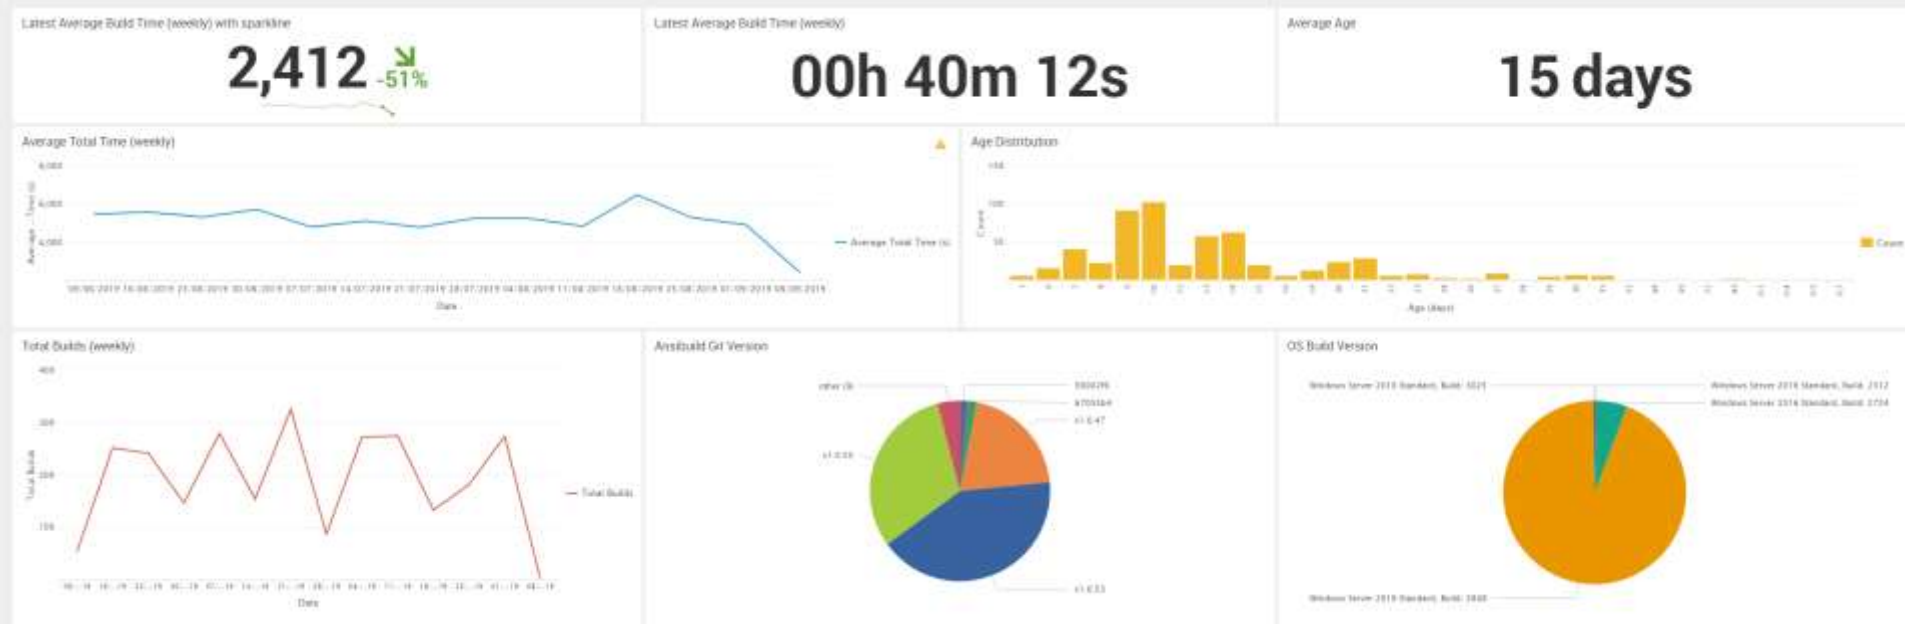
## **Compute Nodes:**

- Splunk forwarder logs inventory information, including image version


## **CMDB:**

- Snapshot of node states are logged to splunk

# Logging - Views



# Logging - Views



# Results

# Results - Deployment




## Result - Builds


	Pre	Post
Teams Involved	5	1
Manual steps	12	2
Build Time (includes test and into 'Farm')	1-2 days	90mins

# Results - Logons

## Server count



## Logons





## Result – Operational

- Servers do not require patching
- Software deployment within 30 days (.net)
- Nodes rebuilt instead of troubleshooting
- Teams can focus on other priorities now build, troubleshooting and management have been reduced.
- Fewer snowflakes !

# Lessons Learnt

# Project Design

- Consider support and external contributors
- Use modular tasks and roles
- Good inventory is critical for simple playbooks

