Splunk

Ravindra Kudache

- *Source of IP traffic*
- *Security threats*

- *System performance*
- *CPU usage & load*


But 1000s of such logs are generated every minute
Check security threats in real-time!!!
Analyze business metrics in real-time!!!

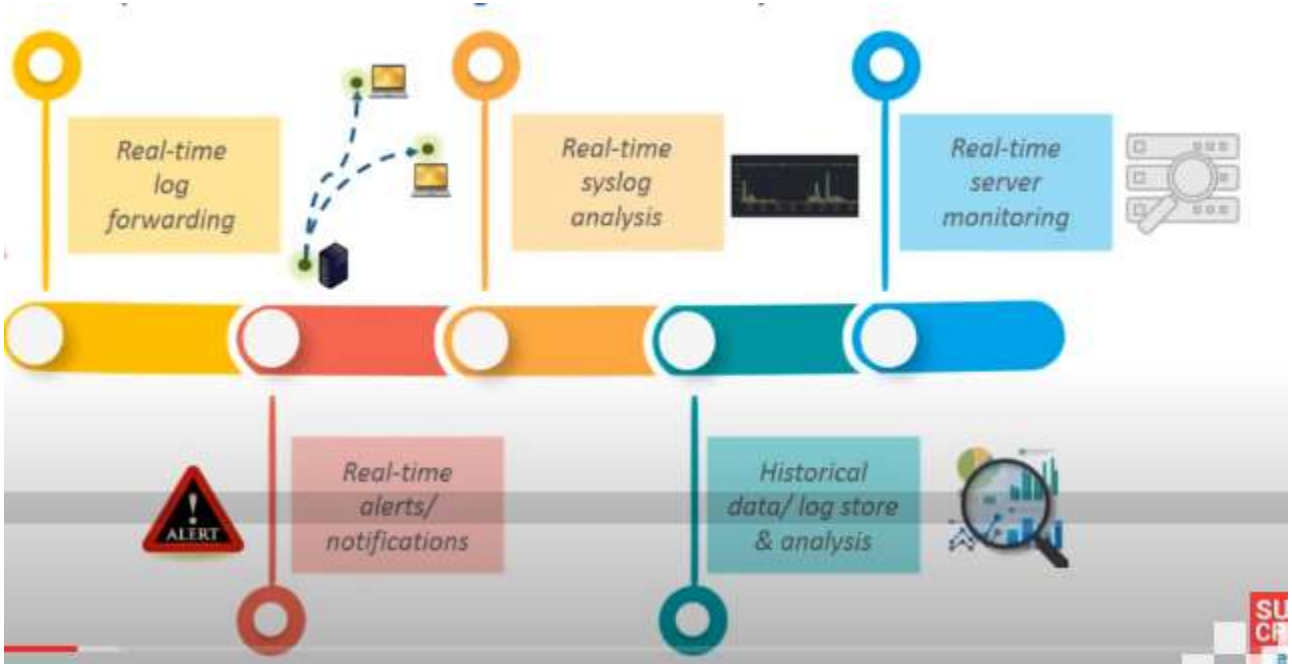
Real-time log forwarding
Real-time syslog analysis
Real-time server monitoring
ALERT
Real-time alerts/ notifications
Historical data/ log store & analysis

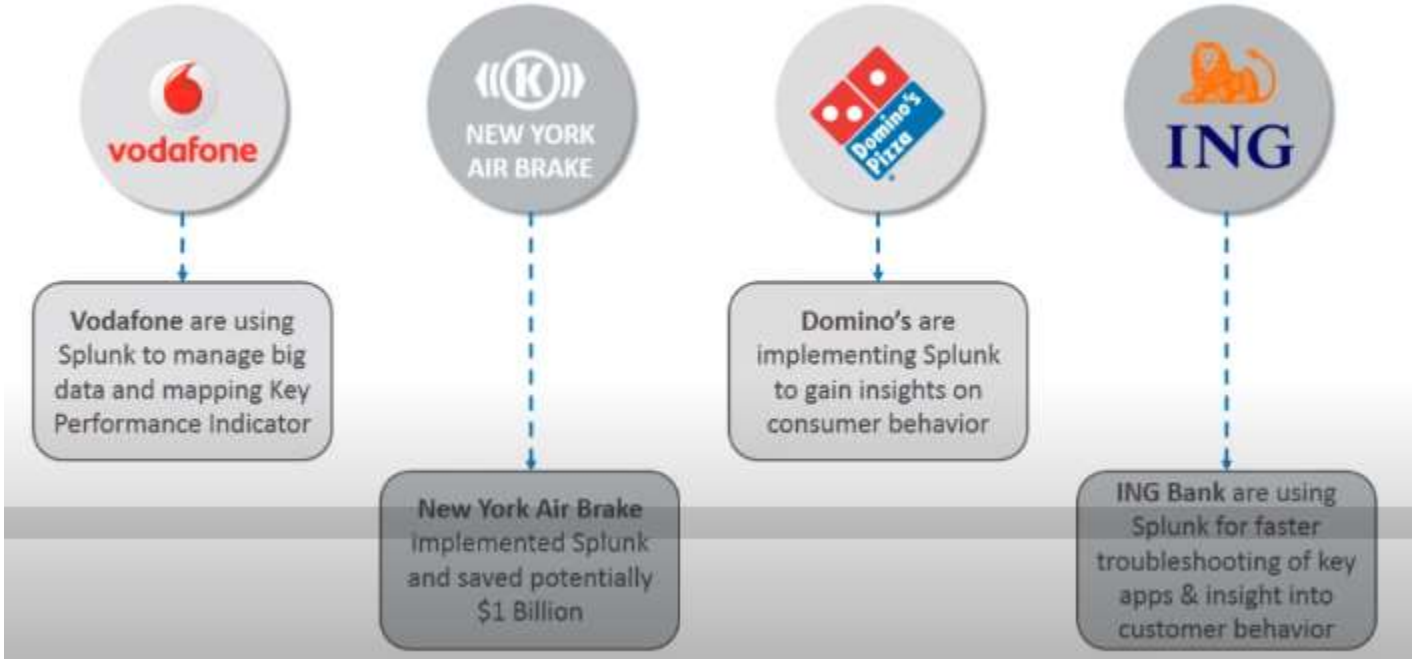
vodafone
Vodafone are using Splunk to manage big data and mapping Key Performance Indicator
NEW YORK AIR BRAKE
New York Air Brake implemented Splunk and saved potentially $1 Billion
Domino's Pizza
Domino's are implementing Splunk to gain insights on consumer behavior
ING
ING Bank are using Splunk for faster troubleshooting of key apps & insight into customer behavior

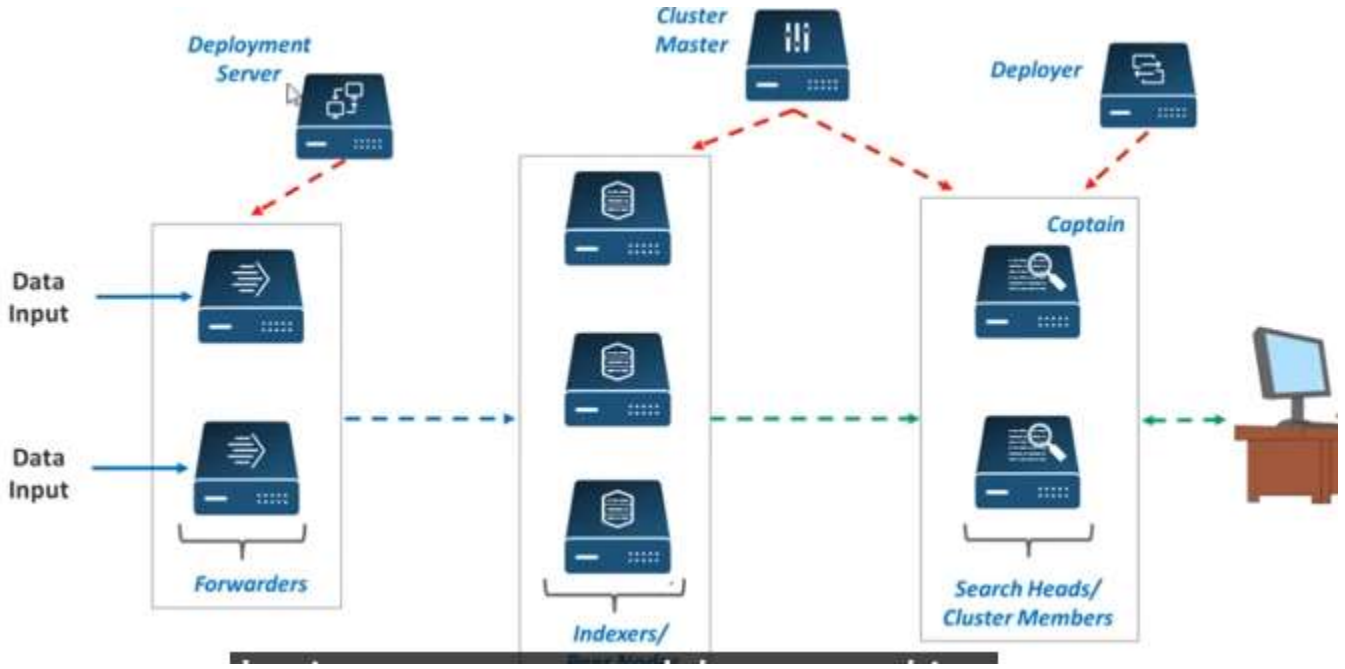
Deployment Server
Cluster Master
Deployer
Data Input
Data Input
Captain
Forwarders
Indexers/
Search Heads/ Cluster Members

Install Splunk

App Dynamic

APM211+-+Introduction+to+AppDynamics+-+Student+Guide.pdf

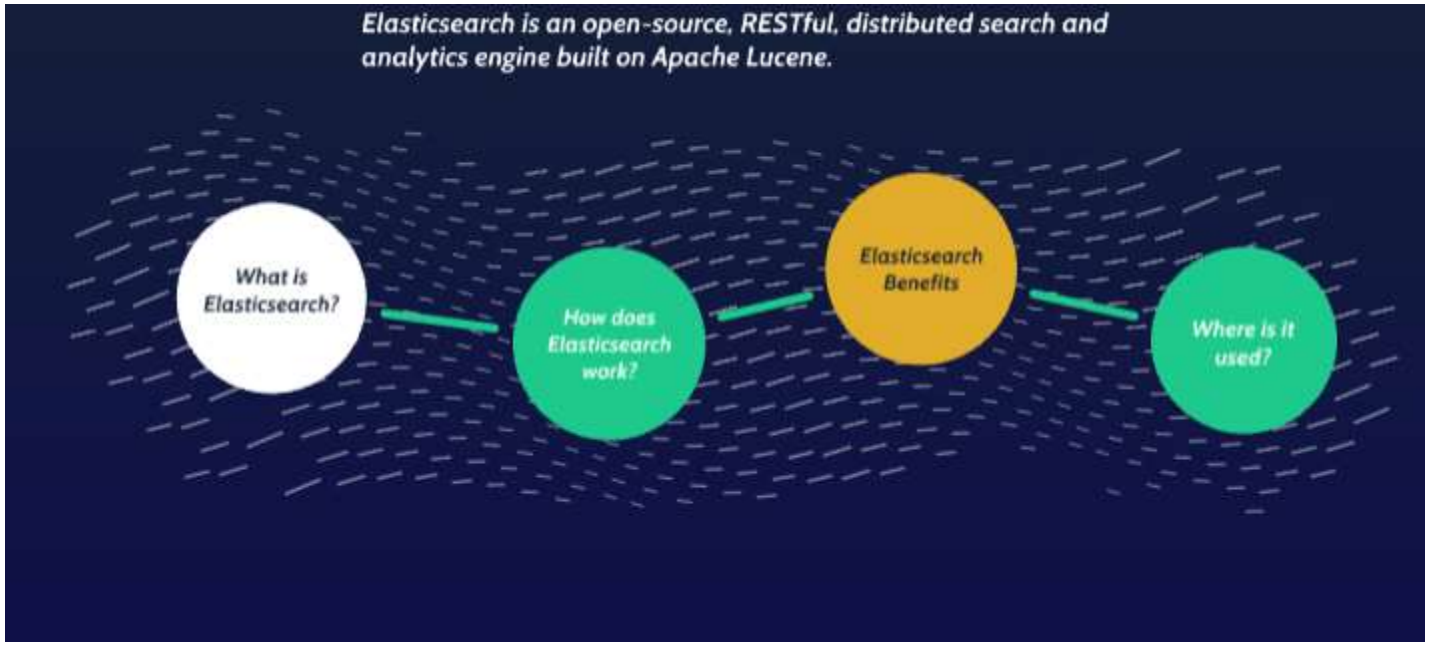
Elasticsearch is an open-source, RESTful, distributed search and analytics engine built on Apache Lucene.
What is Elasticsearch?
How does Elasticsearch work?
Elasticsearch Benefits
Where is it used?


• Elasticsearch was released in 2010, it has quickly become the most popular search engine, and is commonly used for log analytics, full-text search, and operational intelligence use cases.
• When coupled with Kibana, a visualization tool, Elasticsearch can be used to provide near-real time analytics using large volumes of log data. Elasticsearch is also popular because of its easy-to-use search APIs which allow you to easily add powerful search capabilities to your applications.

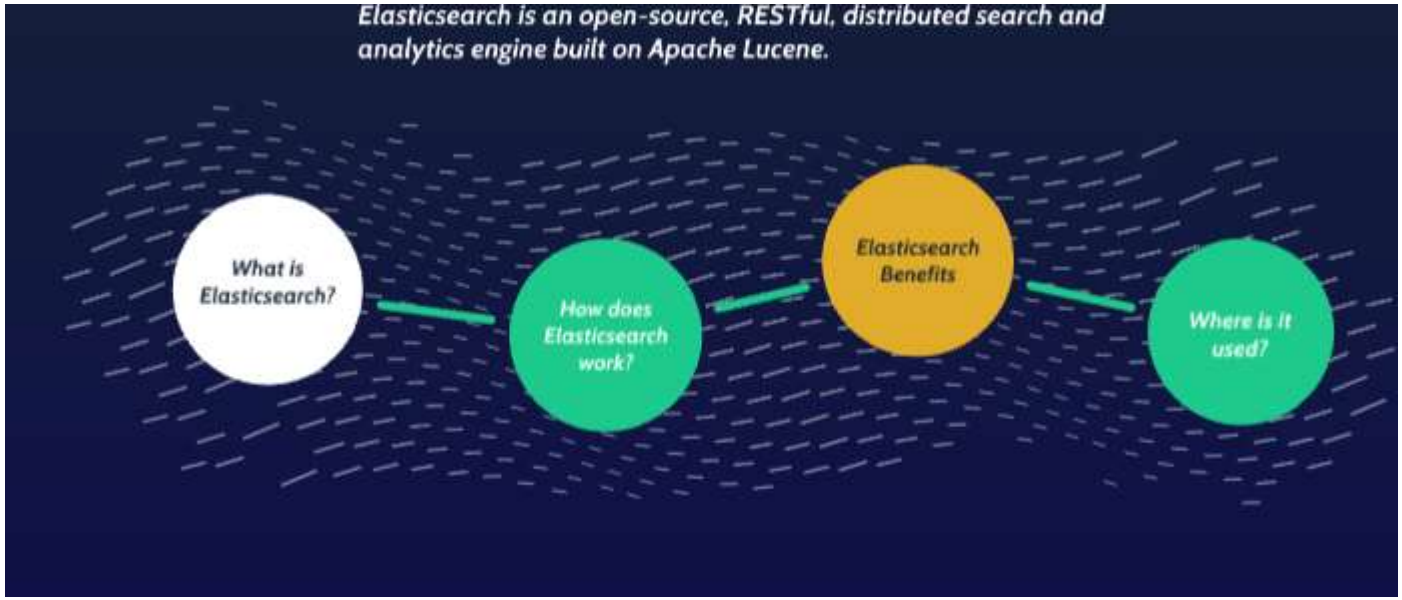
Elasticsearch is an open-source, RESTful, distributed search and analytics engine built on Apache Lucene.
What is Elasticsearch?
How does Elasticsearch work?
Elasticsearch Benefits
Where is it used?

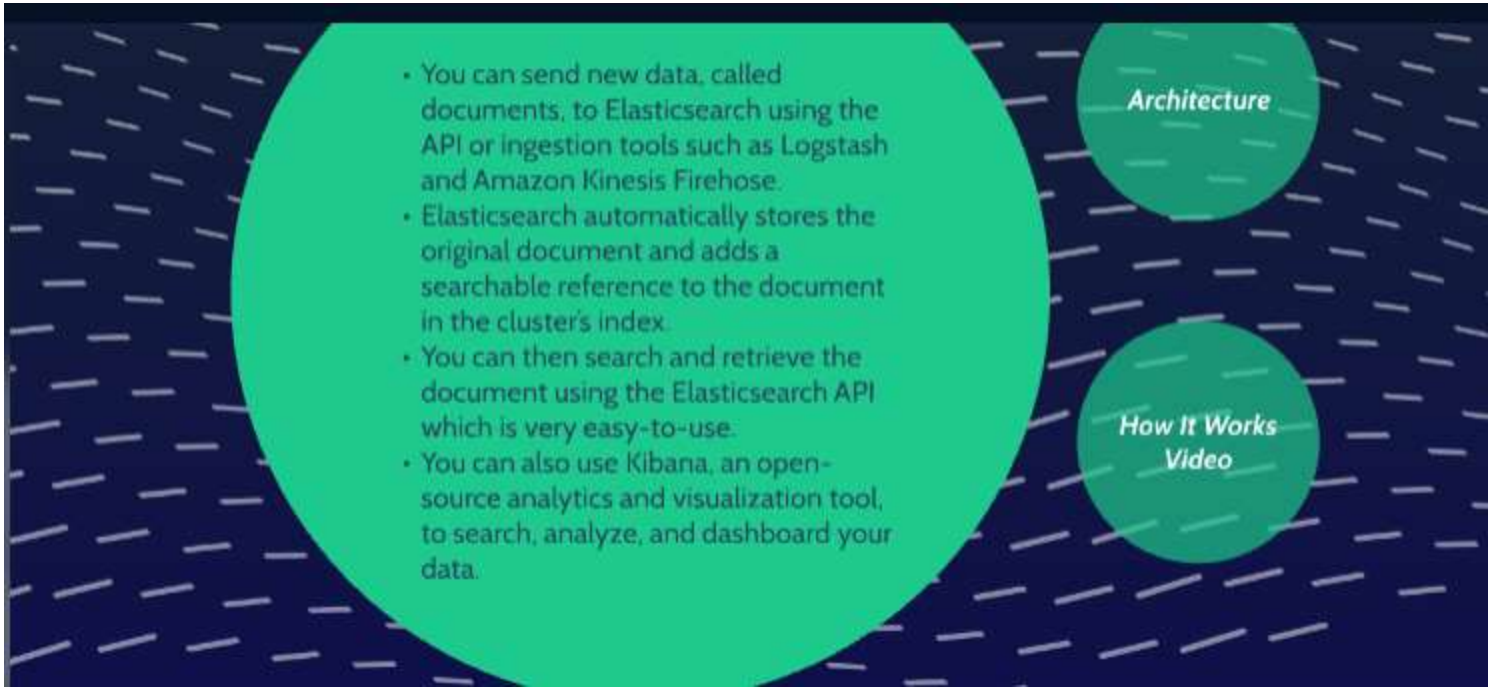
• You can send new data, called documents, to Elasticsearch using the API or ingestion tools such as Logstash and Amazon Kinesis Firehose.
• Elasticsearch automatically stores the original document and adds a searchable reference to the document in the cluster's index.
• You can then search and retrieve the document using the Elasticsearch API which is very easy-to-use.
• You can also use Kibana, an open-source analytics and visualization tool, to search, analyze, and dashboard your data.
Architecture
How It Works Video

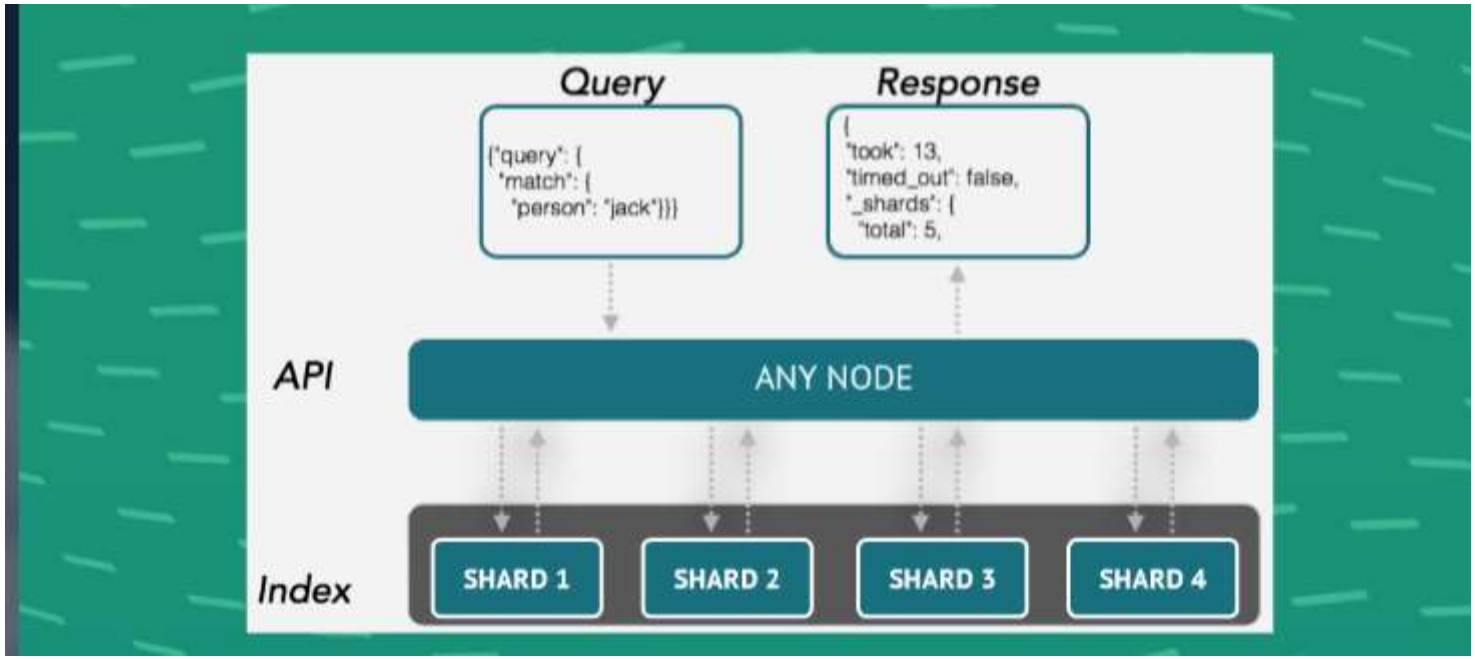
Query
{"query": {
"match": {
"person": "jack"}}}
Response
{
"took": 13,
"timed_out": false,
"_shards": {
"total": 5,
API
ANY NODE
Index
SHARD 1
SHARD 2
SHARD 3
SHARD 4


BENEFITS
QUERY
ANALYZE
SPEED
SCALABILITY
COMPLIMENTARY TOOLING AND PLUG-INS
NEAR REAL-TIME INDEX UPDATES
CLIENT LIBRARIES
MACHINE LEARNING
SECURITY
Speed and Scalability
Complimentary tooling and plug ins
Real-time updates and Client libs
Machine learning and Security

QUERY

Be Curious. Ask Your Data Questions of All Kinds.

- Perform and combine many types of searches – structured, unstructured, geo, metric – any way you want.
- All data types are welcome.

ANALYZE

Step Back and Understand the Bigger Picture.

- It's one thing to find the 10 best documents to match your query. But how do you make sense of, say, a billion log lines?
- Elasticsearch aggregations let you zoom out to explore trends and patterns in your data.

SPEED

Elasticsearch Is Fast. Really, Really Fast.

- When you get answers instantly, your relationship with your data changes. You can afford to iterate and cover more ground.
- And since everything is indexed, you're never left with index envy. You can leverage and access all of your data at ludicrously awesome speeds.

SCALABILITY

Run It on Your Laptop. Or Hundreds of Servers with Petabytes of Data.

- Go from prototype to production seamlessly; you talk to Elasticsearch running on a single node the same way you would in a 300-node cluster.
- It scales horizontally to handle kajillions of events per second, while automatically managing how indices and queries are distributed across the cluster for oh-so smooth operations.

COMPLIMENTARY TOOLING AND PLUG-INS

- Elasticsearch comes integrated with Kibana, a popular visualization and reporting tool.
- It also offers built-in integration with Logstash to easily transform source data using pre-defined templates and load it data into your index.
- In addition, you can use a number of open-source Elasticsearch plug-ins such as language analyzers and suggesters to readily add rich functionality to your applications.
- Put the real-time search and analytics features of Elasticsearch to work on your big data by using the Elasticsearch-Hadoop (ES-Hadoop) connector. It's the best of two worlds colliding.

NEAR REAL-TIME INDEX UPDATES

Elasticsearch Is Fast. Really, Really Fast.

- Elasticsearch index updates such as adding a new document to the index usually take one second or less before the updated data is available for search.
- This lets you use Elasticsearch for near real-time use cases such as application monitoring and anomaly detection.

CLIENT LIBRARIES

Support for your Favorite Development Language

- A variety of open source clients are available for Elasticsearch developers.
- Supported languages include Java, Python, .NET, SQL, PHP, JavaScript, Node.js, Ruby, and many others.

MACHINE LEARNING

It Catches What You Might Miss, All by Itself.

- Complex, fast-moving datasets make it nearly impossible to spot infrastructure problems, intruders, or business issues as they happen using rules or humans looking at dashboards.
- Elastic machine learning features automatically model the behavior of your Elasticsearch data – trends, periodicity, and more – in real time to identify issues faster, streamline root cause analysis, and reduce false positives.

SECURITY

Protect Your Data in the Elastic Stack.

- Elastic Stack security features give the right access to the right people.
- IT, operations, and application teams rely on them to manage well-intentioned users and keep nefarious actors at bay, while executives and customers can rest easy knowing data stored in the Elastic Stack is safe and secure.


Uses of Elasticsearch
• Amazon Web Services
• Adobe Systems
• Center for Open Science
• CERN
• Facebook
• Foursquare
• GitHub
• Lichess
• Mozilla
• Netflix
• Oracle Corporation
• Pixabay
• Quizlet
• Quora
• Reverb
• SeatGeek
• Slurm Workload Manager
• SoundCloud
• Stack Exchange
• StumbleUpon
• Team Foundation Server
• Vimeo
• Wikimedia Foundation
• Zalando SE
SAMPLE CODE SNIPPET
SAMPLE CODE SNIPPET

Java code :

```
Connection conn =
 DriverManager.getConnection("jdbc:elasticsearch:user=myuseraccount;password=mypassword;");

    Statement stat = conn.createStatement();

    ResultSet rs=stat.executeQuery("SELECT * FROM DocumentDB");

    while(rs.next()){
     for(int i=1;i<=rs.getMetaData().getColumnCount();i++)
     {
      System.out.println(rs.getMetaData().getColumnName(i) +"="+rs.getString(i));
     }
    }
```

Java code :

```
RestHighLevelClient client = new RestHighLevelClient(RestClient.builder(
        new HttpHost("localhost", 9200, "http")));

SearchSourceBuilder searchSourceBuilder = new SearchSourceBuilder();

searchSourceBuilder.query(QueryBuilders.matchAllQuery());

searchSourceBuilder.aggregation(AggregationBuilders.terms("top_10_states").field("state").size(10));

SearchRequest searchRequest = new SearchRequest();

searchRequest.indices("social-*");

searchRequest.source(searchSourceBuilder);

SearchResponse searchResponse = client.search(searchRequest);
```